

AbhijitMandal_DSC540_Week3-4Ex

April 4, 2021

0.0.1 DSC 540 Week 3-4

Abhijit Mandal

0.0.2 Activity 5: Generating Statistics from a CSV File

- Load the necessary libraries.
- Read in the Boston housing dataset (given as a .csv file) from the local directory.
- Check the first 10 records. Find the total number of records.
- Create a smaller DataFrame with columns that do not include CHAS, NOX, B, and LSTAT.
- Check the last seven records of the new DataFrame you just created.
- Plot the histograms of all the variables (columns) in the new DataFrame.
- Plot them all at once using a for loop. Try to add a unique title to a plot.
- Create a scatter plot of crime rate versus price.
- Plot using $\log_{10}(\text{crime})$ versus price.
- Calculate some useful statistics, such as mean rooms per dwelling, median age, mean distances to five Boston employment centers, and the percentage of houses with a low price ($< \$20,000$).

0.0.3 Load the necessary libraries.

```
[100]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

0.0.4 Read in the Boston housing dataset (given as a .csv file) from the local directory.

```
[101]: bostonHousingDF=pd.read_csv("../Boston_housing.csv")
```

0.0.5 Check first 10 records

```
[102]: bostonHousingDF.head(10)
```

```
[102]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	\
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3	
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242	17.8	
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	

5	0.02985	0.0	2.18	0	0.458	6.430	58.7	6.0622	3	222	18.7
6	0.08829	12.5	7.87	0	0.524	6.012	66.6	5.5605	5	311	15.2
7	0.14455	12.5	7.87	0	0.524	6.172	96.1	5.9505	5	311	15.2
8	0.21124	12.5	7.87	0	0.524	5.631	100.0	6.0821	5	311	15.2
9	0.17004	12.5	7.87	0	0.524	6.004	85.9	6.5921	5	311	15.2

	B	LSTAT	PRICE
0	396.90	4.98	24.0
1	396.90	9.14	21.6
2	392.83	4.03	34.7
3	394.63	2.94	33.4
4	396.90	5.33	36.2
5	394.12	5.21	28.7
6	395.60	12.43	22.9
7	396.90	19.15	27.1
8	386.63	29.93	16.5
9	386.71	17.10	18.9

0.0.6 Find the total number of records.

```
[103]: bostonHousingDF.shape
```

```
[103]: (506, 14)
```

0.0.7 Create a smaller DataFrame with columns which do not include 'CHAS', 'NOX', 'B', and 'LSTAT'

```
[104]: bostonHousingSmallDF =   
↳ bostonHousingDF[['CRIM', 'ZN', 'INDUS', 'RM', 'AGE', 'DIS', 'RAD', 'TAX', 'PTRATIO', 'PRICE']]
```

0.0.8 Check the last 7 records of the new DataFrame you just created


```
[105]: bostonHousingSmallDF.tail(7)
```

```
[105]:
```


	CRIM	ZN	INDUS	RM	AGE	DIS	RAD	TAX	PTRATIO	PRICE
499	0.17783	0.0	9.69	5.569	73.5	2.3999	6	391	19.2	17.5
500	0.22438	0.0	9.69	6.027	79.7	2.4982	6	391	19.2	16.8
501	0.06263	0.0	11.93	6.593	69.1	2.4786	1	273	21.0	22.4
502	0.04527	0.0	11.93	6.120	76.7	2.2875	1	273	21.0	20.6
503	0.06076	0.0	11.93	6.976	91.0	2.1675	1	273	21.0	23.9
504	0.10959	0.0	11.93	6.794	89.3	2.3889	1	273	21.0	22.0
505	0.04741	0.0	11.93	6.030	80.8	2.5050	1	273	21.0	11.9

0.0.9 Can you plot histograms of all the variables (columns) in the new DataFrame?


```
[106]: for c in bostonHousingSmallDF.columns:  
        plt.title("Plot of "+c,fontsize=15)  
        plt.hist(bostonHousingSmallDF[c],bins=20)  
        plt.show()
```





Plot of ZN




Plot of INDUS







Plot of DIS




Plot of RAD




Plot of TAX




Plot of PTRATIO






0.0.10 Crime rate could be an indicator of house price (people don't want to live in high-crime areas). Create a scatter plot of crime rate vs. Price.

```
[107]: plt.scatter(bostonHousingSmallDF['CRIM'],bostonHousingSmallDF['PRICE'])  
plt.show()
```



0.0.11 We can understand the relationship better if we plot $\log_{10}(\text{crime})$ vs. Price

```
[108]: plt.scatter(np.  
    ↪ log10(bostonHousingSmallDF['CRIM']), bostonHousingSmallDF['PRICE'], c='red')  
plt.title("Crime rate (Log) vs. Price plot", fontsize=18)  
plt.xlabel("Log of Crime rate", fontsize=15)  
plt.ylabel("Price", fontsize=15)  
plt.grid(True)  
plt.show()
```



0.0.12 Calculate the mean rooms per dwelling?

```
[109]: bostonHousingSmallDF['RM'].mean()
```

```
[109]: 6.284634387351788
```

0.0.13 Calculate median Age?

```
[110]: bostonHousingSmallDF['AGE'].median()
```

```
[110]: 77.5
```

0.0.14 Calculate average (mean) distances to five Boston employment centres?

```
[111]: bostonHousingSmallDF['DIS'].mean()
```

```
[111]: 3.795042687747034
```

0.0.15 calculate the percentage of houses with low price (< \$20,000)?

```
[112]: low_price=bostonHousingSmallDF['PRICE']<20
# This creates a Boolean array of True, False
print(low_price)
pcnt=low_price.mean()*100
print("\nPercentage of house with <20,000 price is: ",pcnt)
```

```
0      False
1      False
2      False
3      False
4      False
```

```
...
501     False
502     False
503     False
504     False
505      True
```

```
Name: PRICE, Length: 506, dtype: bool
```

```
Percentage of house with <20,000 price is:  41.50197628458498
```

Activity 6: Working with the Adult Income Dataset (UCI)

- Load the necessary libraries.
- Read the adult income dataset from the following URL:
<https://github.com/TrainingByPackt/Data-Wrangling-with-Python/blob/master/Chapter04/Activity06/>
- Create a script that will read a text file line by line.
- Add a name of Income for the response variable to the dataset.
- Find the missing values.
- Create a DataFrame with only age, education, and occupation by using subsetting.
- Plot a histogram of age with a bin size of 20.
- Create a function to strip the whitespace characters.

- Use the apply method to apply this function to all the columns with string values, create a new column, copy the values from this new column to the old column, and drop the new column.
- Find the number of people who are aged between 30 and 50.
- Group the records based on age and education to find how the mean age is distributed.
- Group by occupation and show the summary statistics of age. Find which profession has the oldest workers on average and which profession has its largest share of the workforce above the 75th percentile.
- Use subset and groupby to find outliers.
- Plot the values on a bar chart.
- Merge the data using common keys.

0.0.16 Read in the adult income data set (given as a .csv file) from the local directory and check first 5 records

```
[113]: df = pd.read_csv("../adult_income_data.csv")
df.head()
```

```
[113]: 39      State-gov    77516  Bachelors    13      Never-married  \
0  50  Self-emp-not-inc    83311  Bachelors    13  Married-civ-spouse
1  38      Private    215646    HS-grad     9      Divorced
2  53      Private    234721     11th     7  Married-civ-spouse
3  28      Private    338409  Bachelors    13  Married-civ-spouse
4  37      Private    284582   Masters    14  Married-civ-spouse

      Adm-clerical  Not-in-family  White  Male  2174  0  40  \
0  Exec-managerial      Husband  White  Male     0  0  13
1  Handlers-cleaners  Not-in-family  White  Male     0  0  40
2  Handlers-cleaners      Husband  Black  Male     0  0  40
3  Prof-specialty      Wife  Black  Female     0  0  40
4  Exec-managerial      Wife  White  Female     0  0  40

      United-States  <=50K
0  United-States  <=50K
1  United-States  <=50K
2  United-States  <=50K
3      Cuba  <=50K
4  United-States  <=50K
```

0.0.17 Create a script that will read a text file line by line and extracts the first line, which is the header of the .csv file:

```
[114]: names = []
with open('../adult_income_names.txt', 'r') as f:
    for line in f:
        f.readline()
        var=line.split(":")[0]
```

```
names.append(var)
names
```

```
[114]: ['age',
        'workclass',
        'fnlwt',
        'education',
        'education-num',
        'marital-status',
        'occupation',
        'relationship',
        'sex',
        'race',
        'capital-gain',
        'capital-loss',
        'hours-per-week',
        'native-country']
```

0.0.18 Add a name (“Income”) for the response variable (last column) to the dataset and read it again with the column names supplied

```
[115]: names.append('Income')
```

```
[116]: df = pd.read_csv("../adult_income_data.csv", names=names)
df.head()
```

```
[116]:
```

	age	workclass	fnlwt	education	education-num	\
0	39	State-gov	77516	Bachelors	13	
1	50	Self-emp-not-inc	83311	Bachelors	13	
2	38	Private	215646	HS-grad	9	
3	53	Private	234721	11th	7	
4	28	Private	338409	Bachelors	13	

	marital-status	occupation	relationship	sex	race	\
0	Never-married	Adm-clerical	Not-in-family	White	Male	
1	Married-civ-spouse	Exec-managerial	Husband	White	Male	
2	Divorced	Handlers-cleaners	Not-in-family	White	Male	
3	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	
4	Married-civ-spouse	Prof-specialty	Wife	Black	Female	

	capital-gain	capital-loss	hours-per-week	native-country	Income
0	2174	0	40	United-States	<=50K
1	0	0	13	United-States	<=50K
2	0	0	40	United-States	<=50K
3	0	0	40	United-States	<=50K
4	0	0	40	Cuba	<=50K

0.0.19 Show a statistical summary of the data set. Did you notice only a small number of columns are included?

```
[117]: df.describe()
```

```
[117]:
```

	age	fnlwgt	education-num	capital-gain	capital-loss \
count	32561.000000	3.256100e+04	32561.000000	32561.000000	32561.000000
mean	38.581647	1.897784e+05	10.080679	1077.648844	87.303830
std	13.640433	1.055500e+05	2.572720	7385.292085	402.960219
min	17.000000	1.228500e+04	1.000000	0.000000	0.000000
25%	28.000000	1.178270e+05	9.000000	0.000000	0.000000
50%	37.000000	1.783560e+05	10.000000	0.000000	0.000000
75%	48.000000	2.370510e+05	12.000000	0.000000	0.000000
max	90.000000	1.484705e+06	16.000000	99999.000000	4356.000000

	hours-per-week
count	32561.000000
mean	40.437456
std	12.347429
min	1.000000
25%	40.000000
50%	40.000000
75%	45.000000
max	99.000000

0.0.20 Many variables in the dataset have multiple factors or classes. Can you write a loop to count and print them?

```
[118]: var_cls = [
    ↪ ['workclass', 'education', 'marital-status', 'occupation', 'relationship', 'race', 'sex', 'native-
for v in var_cls:
    classes=df[v].unique()
    num_classes = df[v].nunique()
    print("There are {} classes in the \"{}\" column. They are: {}".
    ↪format(num_classes,v,classes))
    print("-"*100)
```

```
There are 9 classes in the "workclass" column. They are: [' State-gov' ' Self-
emp-not-inc' ' Private' ' Federal-gov' ' Local-gov'
' ?' ' Self-emp-inc' ' Without-pay' ' Never-worked']
```

```
-----
There are 16 classes in the "education" column. They are: [' Bachelors' ' HS-
grad' ' 11th' ' Masters' ' 9th' ' Some-college'
' Assoc-acdm' ' Assoc-voc' ' 7th-8th' ' Doctorate' ' Prof-school'
' 5th-6th' ' 10th' ' 1st-4th' ' Preschool' ' 12th']
-----
```

There are 7 classes in the "marital-status" column. They are: [' Never-married'
' Married-civ-spouse' ' Divorced'
' Married-spouse-absent' ' Separated' ' Married-AF-spouse' ' Widowed']

There are 15 classes in the "occupation" column. They are: [' Adm-clerical' '
Exec-managerial' ' Handlers-cleaners' ' Prof-specialty'
' Other-service' ' Sales' ' Craft-repair' ' Transport-moving'
' Farming-fishing' ' Machine-op-inspct' ' Tech-support' ' ?'
' Protective-serv' ' Armed-Forces' ' Priv-house-serv']

There are 6 classes in the "relationship" column. They are: [' Not-in-family' '
Husband' ' Wife' ' Own-child' ' Unmarried'
' Other-relative']

There are 2 classes in the "race" column. They are: [' Male' ' Female']

There are 5 classes in the "sex" column. They are: [' White' ' Black' ' Asian-
Pac-Islander' ' Amer-Indian-Eskimo' ' Other']

There are 42 classes in the "native-country" column. They are: [' United-States'
' Cuba' ' Jamaica' ' India' ' ?' ' Mexico' ' South'
' Puerto-Rico' ' Honduras' ' England' ' Canada' ' Germany' ' Iran'
' Philippines' ' Italy' ' Poland' ' Columbia' ' Cambodia' ' Thailand'
' Ecuador' ' Laos' ' Taiwan' ' Haiti' ' Portugal' ' Dominican-Republic'
' El-Salvador' ' France' ' Guatemala' ' China' ' Japan' ' Yugoslavia'
' Peru' ' Outlying-US(Guam-USVI-etc)' ' Scotland' ' Trinidad&Tobago'
' Greece' ' Nicaragua' ' Vietnam' ' Hong' ' Ireland' ' Hungary'
' Holand-Netherlands']

0.0.21 Is there any missing (NULL) data in the dataset? Write a single line of code to show this for all counms

```
[119]: df.isnull().sum()
```

```
[119]: age                0  
workclass              0  
fnlwgt                 0  
education              0  
education-num          0
```

```
marital-status    0
occupation        0
relationship      0
sex              0
race             0
capital-gain      0
capital-loss      0
hours-per-week    0
native-country    0
Income           0
dtype: int64
```

0.0.22 Create a DataFrame with only age, education, and occupation by using sub-setting:


```
[124]: df_subset = df[['age', 'education', 'race', 'occupation']]
df_subset.head()
```

```
[124]:   age  education  race      occupation
0   39  Bachelors  Male    Adm-clerical
1   50  Bachelors  Male    Exec-managerial
2   38   HS-grad  Male  Handlers-cleaners
3   53     11th  Male  Handlers-cleaners
4   28  Bachelors  Female  Prof-specialty
```

0.0.23 Show the histogram of age with bin size = 20

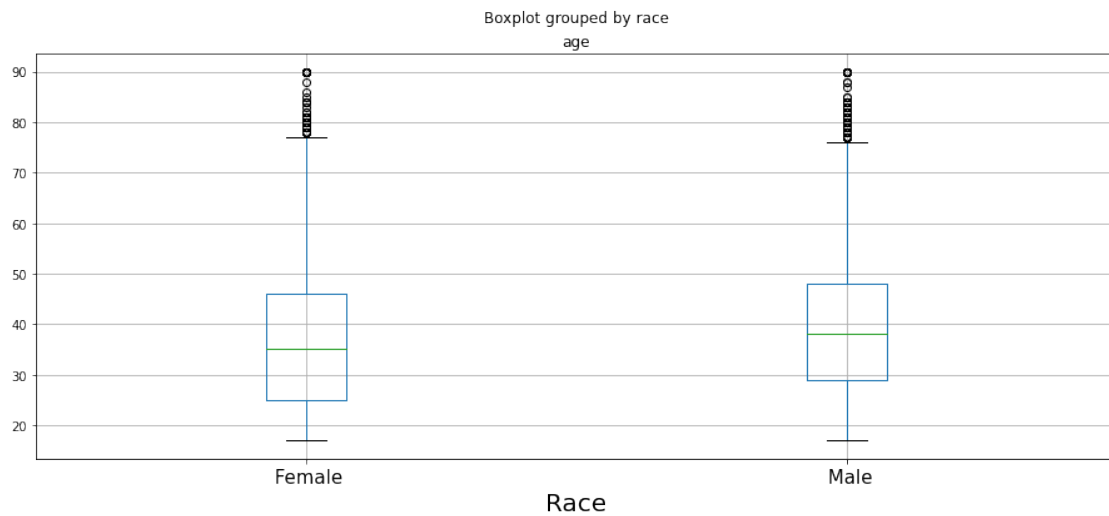
```
[125]: df_subset['age'].hist(bins=20)
```

```
[125]: <AxesSubplot:>
```



0.0.24 Show boxplots of age grouped by race (Use a long figure size 15x6 and make x ticks font size 15)

```
[126]: df_subset.boxplot(column='age',by='race',figsize=(15,6))
plt.xticks(fontsize=15)
plt.xlabel("Race",fontsize=20)
plt.show()
```



0.0.25 Create a function to strip the whitespace characters

```
[127]: def strip_whitespace(s):  
        return s.strip()
```

0.0.26 Use the 'apply' method to apply this function to all the columns with string values, create a new column, copy the values from this new column to the old column, and drop the new column.

```
[128]: df_subset['education_stripped']=df['education'].apply(strip_whitespace)  
df_subset['education']=df_subset['education_stripped']  
df_subset.drop(labels=['education_stripped'],axis=1,inplace=True)  
  
df_subset['occupation_stripped']=df['occupation'].apply(strip_whitespace)  
df_subset['occupation']=df_subset['occupation_stripped']  
df_subset.drop(labels=['occupation_stripped'],axis=1,inplace=True)  
  
df_subset['race_stripped']=df['race'].apply(strip_whitespace)  
df_subset['race']=df_subset['race_stripped']  
df_subset.drop(labels=['race_stripped'],axis=1,inplace=True)
```

```
<ipython-input-128-5e98db29da69>:1: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy  
df_subset['education_stripped']=df['education'].apply(strip_whitespace)  
<ipython-input-128-5e98db29da69>:2: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy  
df_subset['education']=df_subset['education_stripped']  
/Users/abhijitmandal/opt/anaconda3/lib/python3.8/site-packages/pandas/core/frame.py:4163: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy  
return super().drop(  
<ipython-input-128-5e98db29da69>:6: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
df_subset['occupation_stripped']=df['occupation'].apply(strip_whitespace)
<ipython-input-128-5e98db29da69>:7: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
df_subset['occupation']=df_subset['occupation_stripped']
<ipython-input-128-5e98db29da69>:11: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
df_subset['race_stripped']=df['race'].apply(strip_whitespace)
<ipython-input-128-5e98db29da69>:12: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
df_subset['race']=df_subset['race_stripped']

0.0.27 Find the number of people who are aged between 30 and 50 (inclusive) by using the following command

```
[129]: df_filtered=df_subset[(df_subset['age']>=30) & (df_subset['age']<=50)]
```

```
[130]: df_filtered.head()
```

```
[130]:
```

	age	education	race	occupation
0	39	Bachelors	Male	Adm-clerical
1	50	Bachelors	Male	Exec-managerial
2	38	HS-grad	Male	Handlers-cleaners
5	37	Masters	Female	Exec-managerial
6	49	9th	Female	Other-service

0.0.28 Find the shape of the filtered DataFrame and specify the index of the tuple as 0 to return the first element

```
[131]: data1=df_filtered.shape[0]  
data1
```

```
[131]: 16390
```

0.0.29 Print the number of black people aged between 30 and 50 using the following command

```
[132]: print("There are {} people of age between 30 and 50 in this dataset.".
        ↪format(data1))
```

There are 16390 people of age between 30 and 50 in this dataset.

0.0.30 Group the records based on occupation to find how the mean age is distributed:

```
[133]: df_subset.groupby('occupation').describe()['age']
```

```
[133]:
```

	count	mean	std	min	25%	50%	75%	max
occupation								
?	1843.0	40.882800	20.336350	17.0	21.0	35.0	61.0	90.0
Adm-clerical	3770.0	36.964456	13.362998	17.0	26.0	35.0	46.0	90.0
Armed-Forces	9.0	30.222222	8.089774	23.0	24.0	29.0	34.0	46.0
Craft-repair	4099.0	39.031471	11.606436	17.0	30.0	38.0	47.0	90.0
Exec-managerial	4066.0	42.169208	11.974548	17.0	33.0	41.0	50.0	90.0
Farming-fishing	994.0	41.211268	15.070283	17.0	29.0	39.0	52.0	90.0
Handlers-cleaners	1370.0	32.165693	12.372635	17.0	23.0	29.0	39.0	90.0
Machine-op-inspct	2002.0	37.715285	12.068266	17.0	28.0	36.0	46.0	90.0
Other-service	3295.0	34.949621	14.521508	17.0	22.0	32.0	45.0	90.0
Priv-house-serv	149.0	41.724832	18.633688	17.0	24.0	40.0	57.0	81.0
Prof-specialty	4140.0	40.517633	12.016676	17.0	31.0	40.0	48.0	90.0
Protective-serv	649.0	38.953775	12.822062	17.0	29.0	36.0	47.0	90.0
Sales	3650.0	37.353973	14.186352	17.0	25.0	35.0	47.0	90.0
Tech-support	928.0	37.022629	11.316594	17.0	28.0	36.0	44.0	73.0
Transport-moving	1597.0	40.197871	12.450792	17.0	30.0	39.0	49.0	90.0

0.0.31 Group by occupation and show the summary statistics of age. Find which profession has the oldest workers on average and which profession has its largest share of workforce above the 75th percentile:

```
[134]: df_subset.groupby('occupation').describe()['age']
```

```
[134]:
```

	count	mean	std	min	25%	50%	75%	max
occupation								
?	1843.0	40.882800	20.336350	17.0	21.0	35.0	61.0	90.0
Adm-clerical	3770.0	36.964456	13.362998	17.0	26.0	35.0	46.0	90.0
Armed-Forces	9.0	30.222222	8.089774	23.0	24.0	29.0	34.0	46.0
Craft-repair	4099.0	39.031471	11.606436	17.0	30.0	38.0	47.0	90.0
Exec-managerial	4066.0	42.169208	11.974548	17.0	33.0	41.0	50.0	90.0
Farming-fishing	994.0	41.211268	15.070283	17.0	29.0	39.0	52.0	90.0
Handlers-cleaners	1370.0	32.165693	12.372635	17.0	23.0	29.0	39.0	90.0
Machine-op-inspct	2002.0	37.715285	12.068266	17.0	28.0	36.0	46.0	90.0
Other-service	3295.0	34.949621	14.521508	17.0	22.0	32.0	45.0	90.0


Priv-house-serv	149.0	41.724832	18.633688	17.0	24.0	40.0	57.0	81.0
Prof-specialty	4140.0	40.517633	12.016676	17.0	31.0	40.0	48.0	90.0
Protective-serv	649.0	38.953775	12.822062	17.0	29.0	36.0	47.0	90.0
Sales	3650.0	37.353973	14.186352	17.0	25.0	35.0	47.0	90.0
Tech-support	928.0	37.022629	11.316594	17.0	28.0	36.0	44.0	73.0
Transport-moving	1597.0	40.197871	12.450792	17.0	30.0	39.0	49.0	90.0

0.0.32 Use subset and groupby to find the outliers

```
[135]: occupation_stats= df_subset.groupby( 'occupation').describe()['age']
```

0.0.33 Plot the values on a bar chart

```
[136]: plt.figure(figsize=(15,8))
plt.barh(y=occupation_stats.index, width=occupation_stats['count'])
plt.yticks(fontsize=13)
plt.show()
```



0.0.34 Practice Merging by common keys: Suppose you are given two datasets where the common key is occupation. Can you merge them?¶

```
[137]: df_1 = df[['age', 'workclass', 'occupation']].sample(5, random_state=101)
df_1.head()
```

```
[137]:
```

	age	workclass	occupation
22357	51	Private	Machine-op-inspct
26009	19	Private	Sales
20734	40	Private	Exec-managerial

```
17695    17    Private    Handlers-cleaners
27908    61    Private          Craft-repair
```

```
[138]: df_2 = df[['education', 'race', 'occupation']].sample(5, random_state=101)
df_2.head()
```

```
[138]:      education    race    occupation
22357    HS-grad    Female    Machine-op-inspct
26009      11th      Male          Sales
20734    HS-grad      Male    Exec-managerial
17695     10th      Male    Handlers-cleaners
27908    7th-8th      Male    Craft-repair
```

```
[139]: df_merged = pd.merge(df_1, df_2, on='occupation', how='inner').drop_duplicates()
df_merged
```

```
[139]:   age workclass    occupation education    race
0    51    Private    Machine-op-inspct    HS-grad    Female
1    19    Private          Sales      11th      Male
2    40    Private    Exec-managerial    HS-grad      Male
3    17    Private    Handlers-cleaners     10th      Male
4    61    Private    Craft-repair     7th-8th      Male
```

0.0.35 Create a series and practice basic arithmetic steps

- Series 1 = 7.3, -2.5, 3.4, 1.5
 - i. Index = 'a', 'c', 'd', 'e'
- Series 2 = -2.1, 3.6, -1.5, 4, 3.1
 - i. Index = 'a', 'c', 'e', 'f', 'g'
- Add Series 1 and Series 2 together and print the results
- Subtract Series 1 from Series 2 and print the results

```
[142]: # series 1
data1 = np.array([7.3, -2.5, 3.4, 1.5])

# providing an index
series1 = pd.Series(data1, index=['a', 'c', 'd', 'e'])
print(series1)
```

```
a      7.3
c     -2.5
d      3.4
e      1.5
dtype: float64
```

```
[143]: # series 2
data2 = np.array([-2.1, 3.6, -1.5, 4, 3.1])
```

```
# providing an index
series2 = pd.Series(data2, index=['a', 'c', 'e', 'f','g'])
print(series2)
```

```
a    -2.1
c     3.6
e    -1.5
f     4.0
g     3.1
dtype: float64
```

```
[145]: # series 1 + series 2
series3 = series1.add(series2, fill_value=10);
print(series3)
```

```
a     5.2
c     1.1
d    13.4
e     0.0
f    14.0
g    13.1
dtype: float64
```

```
[147]: # series 1 - series 2
series4 = series1.subtract(series2, fill_value=10);
print(series4)
```

```
a     9.4
c    -6.1
d    -6.6
e     3.0
f     6.0
g     6.9
dtype: float64
```

```
[ ]:
```